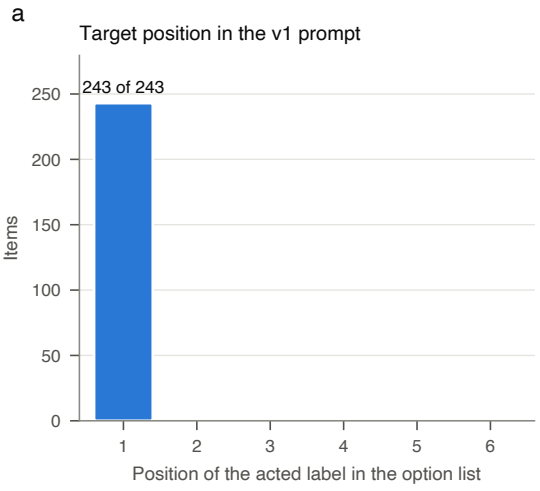
a
Target position in the v1 prompt
243 of 243
Items
250
200
150
100
50
0
1
2
3
4
5
6
Position of the acted label in the option list

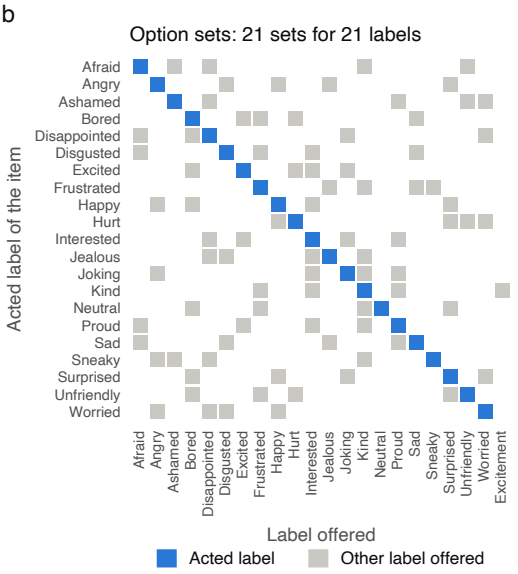
b
Option sets: 21 sets for 21 labels
Acted label of the item
Afraid
Angry
Ashamed
Bored
Disappointed
Disgusted
Excited
Frustrated
Happy
Hurt
Interested
Jealous
Joking
Kind
Neutral
Proud
Sad
Sneaky
Surprised
Unfriendly
Worried
Excitement
Label offered
Acted label
Other label offered

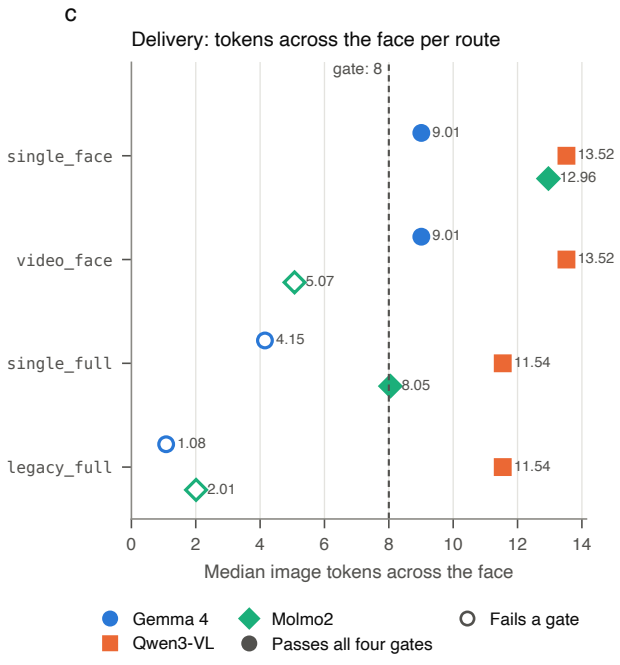
c
Delivery: tokens across the face per route
gate: 8
single_face
video_face
single_full
legacy_full
9.01
13.52
12.96
9.01
13.52
5.07
4.15
11.54
8.05
1.08
11.54
2.01
0
2
4
6
8
10
12
14
Median image tokens across the face
Gemma 4
Qwen3-VL
Molmo2
Passes all four gates
Fails a gate

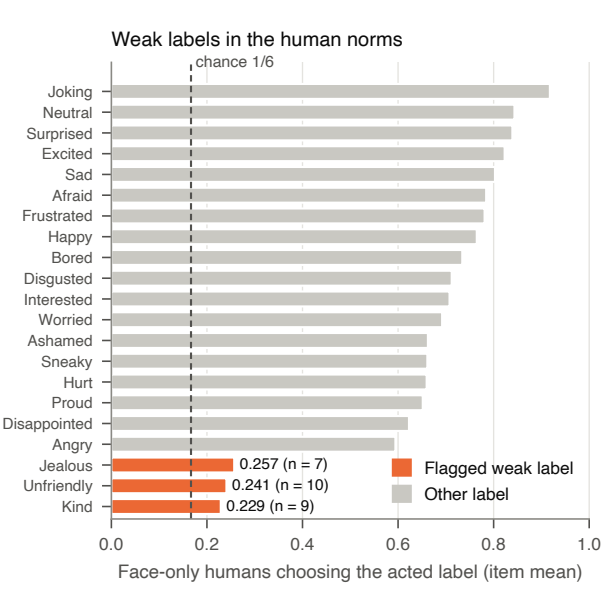
d
Weak labels in the human norms
chance 1/6
Joking
Neutral
Surprised
Excited
Sad
Afraid
Frustrated
Happy
Bored
Disgusted
Interested
Worried
Ashamed
Sneaky
Hurt
Proud
Disappointed
Angry
Jealous
Unfriendly
Kind
0.257 (n = 7)
0.241 (n = 10)
0.229 (n = 9)
Flagged weak label
Other label
0.0
0.2
0.4
0.6
0.8
1.0
Face-only humans choosing the acted label (item mean)

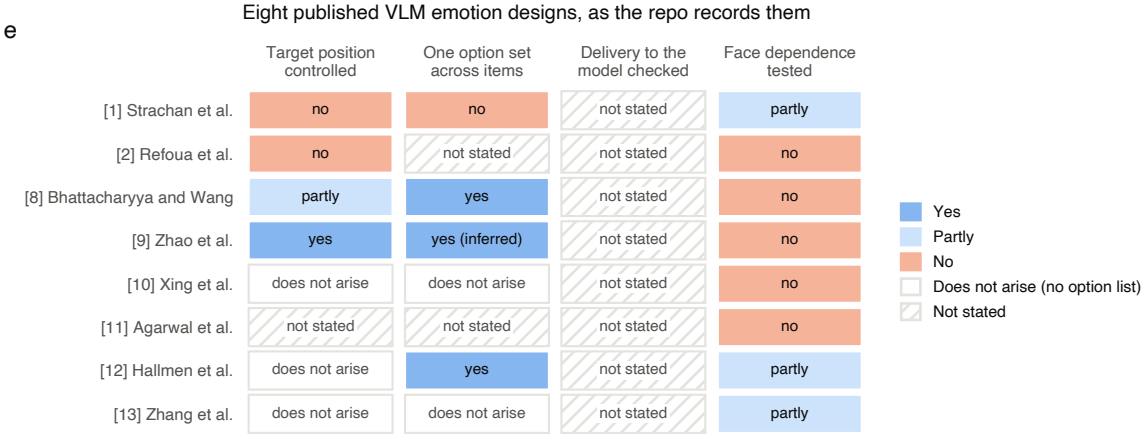
e
Eight published VLM emotion designs, as the repo records them
Target position controlled
One option set across items
Delivery to the model checked
Face dependence tested
[1] Strachan et al.
no
no
not stated
partly
[2] Refoua et al.
no
not stated
not stated
no
[8] Bhattacharyya and Wang
partly
yes
not stated
no
[9] Zhao et al.
yes
yes (inferred)
not stated
no
[10] Xing et al.
does not arise
does not arise
not stated
no
[11] Agarwal et al.
not stated
not stated
not stated
no
[12] Hallmen et al.
does not arise
yes
not stated
partly
[13] Zhang et al.
does not arise
does not arise
not stated
partly
Yes
Partly
No
Does not arise (no option list)
Not stated